

# K8s实践作业

本次事件作业采用的是阿里云上的云服务器，云服务器的配置如下图。



## 1.设置主机名

```
1 | sudo hostnamectl set-hostname k8s-master
2 | sudo hostnamectl set-hostname k8s-worker
```

可以通过hostname查看是否设置正确

## 2.修改本地DNS，让两台机器之间可以相互ping通

```
1 | sudo nano /etc/hosts
```

在其中添加这两行

```
172.24.153.39 k8s-master k8s-master
172.23.193.146 k8s-worker k8s-worker
```

尝试是否能够ping通

```
1 | root@k8s-worker:~# ping k8s-master
2 | PING k8s-master (172.24.153.39) 56(84) bytes of data.
3 | 64 bytes from k8s-master (172.24.153.39): icmp_seq=1 ttl=64 time=2.38 ms
4 | 64 bytes from k8s-master (172.24.153.39): icmp_seq=2 ttl=64 time=2.27 ms
5 |
6 | root@k8s-master:~# ping k8s-worker
7 | PING k8s-worker (172.23.193.146) 56(84) bytes of data.
8 | 64 bytes from k8s-worker (172.23.193.146): icmp_seq=1 ttl=64 time=2.76 ms
9 | 64 bytes from k8s-worker (172.23.193.146): icmp_seq=2 ttl=64 time=2.62 ms
10 |
```

出现如上即代表可以ping通。

## 3.配置内核路由转发及网桥过滤

```
1 | cat <<EOF | sudo tee /etc/modules-load.d/k8s.conf
2 | overlay
3 | br_netfilter
4 | EOF
5 | sudo modprobe overlay
6 | sudo modprobe br_netfilter
7 | # 设置所需的 sysctl 参数，参数在重新启动后保持不变
```

```

8 cat > /etc/sysctl.d/k8s.conf << EOF
9 net.bridge.bridge-nf-call-ip6tables = 1
10 net.bridge.bridge-nf-call-iptables = 1
11 net.ipv4.ip_forward = 1
12 EOF
13 # 查看是否加载
14 # lsmod | grep br_netfilter
15 # br_netfilter          22256  0
16 # bridge                151336  1 br_netfilter
17 # 应用 sysctl 参数而不重新启动
18 sudo sysctl --system

```

通过运行以下指令确认 `net.bridge.bridge-nf-call-iptables`、`net.bridge.bridge-nf-call-ip6tables` 和 `net.ipv4.ip_forward` 系统变量在你的 `sysctl` 配置中被设置为 1:

```

1 root@k8s-master:/etc/containerd# sysctl net.bridge.bridge-nf-call-iptables
net.bridge.bridge-nf-call-ip6tables net.ipv4.ip_forward
2 net.bridge.bridge-nf-call-iptables = 1
3 net.bridge.bridge-nf-call-ip6tables = 1
4 net.ipv4.ip_forward = 1

```

代表设置成功

## 4. 下载容器运行时

我们这里选择下载 `containerd` 作为我们的容器运行时，具体下载方案参考 `containerd` 的官方文档，如下：这里简单就以我们的 `k8s-master` 节点举例

```

1 root@k8s-master:/home# ls
2 cni-plugins-linux-amd64-v1.4.1.tgz  containerd-1.7.16-linux-amd64.tar.gz
   runc.amd64

```

将安装必要的压缩包文件传入服务器的 `/home` 文件夹下。

安装 `containerd`

```

1 root@k8s-master:/home# tar Cxzf /usr/local containerd-1.7.16-linux-
amd64.tar.gz
2 bin/
3 bin/containerd-shim-runc-v2
4 bin/containerd-stress
5 bin/containerd
6 bin/containerd-shim-runc-v1
7 bin/ctr
8 bin/containerd-shim

```

然后我们将 `containerd.service` 这个文件从 <https://raw.githubusercontent.com/containerd/containerd/main/containerd.service> 拷贝到服务器的 `/usr/lib/systemd/system/containerd.service` 然后运行下面指令

```

1 systemctl daemon-reload
2 systemctl enable --now containerd

```

安装 `runc`

```
1 root@k8s-master:/home# install -m 755 runc.amd64 /usr/local/sbin/runc
```

## 安装 cni

```
1 root@k8s-master:/home# mkdir -p /opt/cni/bin
2 root@k8s-master:/home# tar Cxzvf /opt/cni/bin cni-plugins-linux-amd64-
  v1.4.1.tgz
3 ./
4 ./LICENSE
5 ./host-device
6 ./dummy
7 ./README.md
8 ./firewall
9 ./macvlan
10 ./bridge
11 ./dhcp
12 ./bandwidth
13 ./tuning
14 ./vlan
15 ./ipvlan
16 ./ptp
17 ./static
18 ./loopback
19 ./tap
20 ./host-local
21 ./sbr
22 ./portmap
23 ./vrf
```

## 生成 containerd 的配置

containerd 默认配置文件在 `/etc/containerd` 目录下, 名称为 `config.toml`。

可以通过如下命令生成默认配置:

```
1 root@k8s-master:/etc# mkdir -p /etc/containerd
2 root@k8s-master:/etc# containerd config default > /etc/containerd/config.toml
```

注意需要修改下面两个地方:

```
1 [plugins."io.containerd.grpc.v1.cri".containerd.runtimes.runc]
2   ...
3   [plugins."io.containerd.grpc.v1.cri".containerd.runtimes.runc.options]
4     SystemdCgroup = true
```

SystemdCgroup 这里需要设为true

另外需要设置沙盒镜像为阿里云。

```
sandbox_image = "registry.aliyuncs.com/google_containers/pause:3.9"
```

设置完之后重启 containerd

```
1 sudo systemctl restart containerd
```

## 5. 下载 kubectl kubeadm kubect1

这里我们参考阿里云上提供的官方下载方法，设置阿里云镜像加速安装。

```
1 # 1. 更新软件源并安装基础依赖
2 sudo apt-get update && sudo apt-get install -y apt-transport-https ca-
  certificates curl
3
4 # 2. 添加Kubernetes GPG密钥（通过阿里云镜像）
5 sudo mkdir -p /etc/apt/keyrings
6 curl -fsSL https://mirrors.aliyun.com/kubernetes-
  new/core/stable/v1.28/deb/Release.key | sudo gpg --dearmor -o
  /etc/apt/keyrings/kubernetes-apt-keyring.gpg
7
8 # 3. 添加Kubernetes APT源（阿里云镜像）
9 echo "deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]
  https://mirrors.aliyun.com/kubernetes-new/core/stable/v1.28/deb/ /" | sudo
  tee /etc/apt/sources.list.d/kubernetes.list
10
11 # 4. 更新软件源并安装组件
12 sudo apt-get update
13 sudo apt-get install -y kubectl kubeadm kubect1
14 sudo apt-mark hold kubectl kubeadm kubect1 # 防止自动升级
```

## 6. 然后我们调用如下指令初始化master节点

```
1 # 使用阿里云镜像源初始化Kubernetes集群 (v1.28.9)
2 sudo kubeadm init \
3   --kubernetes-version=v1.28.9 \
4   --pod-network-cidr=10.244.0.0/16 \
5   --image-repository=registry.aliyuncs.com/google_containers
```

最后可以看到终端有下面这行指令

```
1 kubeadm join 172.24.153.39:6443 --token yb3v37.o8p3emy5aie3yro4 \
2   --discovery-token-ca-cert-hash
  sha256:cc3365e7c219893c51ca3c3278a41cf762339100c209c6b8c10a0de213300b0a
```

同时我们看到下面的提示

```
1 To start using your cluster, you need to run the following as a regular user:
2
3   mkdir -p $HOME/.kube
4   sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
5   sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

执行要求的上述三条指令。

这个时候我们看到

```
1 root@k8s-master:/home# kubectl get node
2 NAME           STATUS    ROLES          AGE   VERSION
3 k8s-master     NotReady control-plane  36m   v1.28.9
```

还没有ready，需要安装网络插件，这里我们选择安装flannel插件

从<https://github.com/flannel-io/flannel/blob/master/Documentation/kube-flannel.yml>处下载 kube-flannel.yml

```
1 #安装插件
2 kubectl create -f kube-flannel.yml
```

一段时间过后调用 `kubectl get node`

```
1 root@k8s-master:/home# kubectl get node
2 NAME          STATUS    ROLES          AGE   VERSION
3 k8s-master    Ready    control-plane  37m   v1.28.9
```

master节点已经就绪了。

## 7.将worker节点也加入集群

调用我们上面的指令

```
1 root@k8s-worker:/etc/containerd# kubeadm join 172.24.153.39:6443 \
2   --token yb3v37.o8p3emy5aie3yro4 \
3   --discovery-token-ca-cert-hash
sha256:cc3365e7c219893c51ca3c3278a41cf762339100c209c6b8c10a0de213300b0asha25
6:cc3365e7c219893c51ca3c3278a41cf762339100c209c6b8c10a0de213300b0a
4 accepts at most 1 arg(s), received 2
5 To see the stack trace of this error execute with --v=5 or higher
6 root@k8s-worker:/etc/containerd# kubeadm join 172.24.153.39:6443 \
7   --token yb3v37.o8p3emy5aie3yro4 \
8   --discovery-token-ca-cert-hash
sha256:cc3365e7c219893c51ca3c3278a41cf762339100c209c6b8c10a0de213300b0asha25
6:cc3365e7c219893c51ca3c3278
9 accepts at most 1 arg(s), received 2
10 To see the stack trace of this error execute with --v=5 or higher
11 root@k8s-worker:/etc/containerd# kubeadm join 172.24.153.39:6443 \
12   --token yb3v37.o8p3emy5aie3yro4 \
13   --discovery-token-ca-cert-hash
sha256:cc3365e7c219893c51ca3c3278a41cf762339100c209c6b8c10a0de213300b0asha25
6:cc3365e7c219893c51ca3c3278
14 [preflight] Running pre-flight checks
15 [preflight] Reading configuration from the cluster...
16 [preflight] FYI: You can look at this config file with 'kubectl -n kube-
system get cm kubeadm-config -o yaml'
17 [kubelet-start] writing kubelet configuration to file
"/var/lib/kubelet/config.yaml"
18 [kubelet-start] writing kubelet environment file with flags to file
"/var/lib/kubelet/kubeadm-flags.env"
19 [kubelet-start] Starting the kubelet
20 [kubelet-start] waiting for the kubelet to perform the TLS Bootstrap...
21
22 This node has joined the cluster:
23 * Certificate signing request was sent to apiserver and a response was
received.
24 * The Kubelet was informed of the new secure connection details.
25
```

```
26 | Run 'kubectl get nodes' on the control-plane to see this node join the cluster.
```

成功加入集群。

一段时间后，worker节点也就绪了。

```
1 | root@k8s-master:/home# kubectl get node
2 | NAME                STATUS    ROLES                  AGE     VERSION
3 | k8s-master           Ready     control-plane          44m     v1.28.9
4 | k8s-worker           Ready     <none>                  2m      v1.28.9
```

添加master标签

```
1 | root@k8s-master:/home# kubectl label nodes k8s-master node-
  role.kubernetes.io/master=
2 | node/k8s-master labeled
3 | root@k8s-master:/home# kubectl get node
4 | NAME                STATUS    ROLES                  AGE     VERSION
5 | k8s-master           Ready     control-plane,master   47m     v1.28.9
6 | k8s-worker           Ready     <none>                  4m29s   v1.28.9
```

这样就跟作业文档里要求的一致了。

## Q1: 请记录所有安装步骤的指令，并简要描述其含义

回答：具体的详实安装教程都已经在了上面了。

## Q2:在两个节点上分别使用 ps aux | grep kube 列出所有和k8s相关的进程，记录其输出，并简要说明各个进程的作用

回答：在master节点上输出如下：

```
root@k8s-master:/home# ps aux | grep kube
root    5047  0.3  2.7 1344076 186608 ?        Ssl  11:46   0:19 kube-controller-manager --allocate-node-cidrs=true --authentication-kubeconfig=/etc/kubernetes/controller-manager.conf --authorization-kubeconfig=/etc/kubernetes/controller-manager.conf --bind-address=127.0.0.1 --client-ca-file=/etc/kubernetes/pki/ca.crt --cluster-cidr=10.244.0.0/16 --cluster-name=kubernetes --cluster-signing-cert-file=/etc/kubernetes/pki/ca.crt --cluster-signing-key-file=/etc/kubernetes/pki/ca.key --controllers=*,bootstrapsigner,tokencleaner --kubeconfig=/etc/kubernetes/controller-manager.conf --leader-elect=true --requestheader-client-ca-file=/etc/kubernetes/pki/front-proxy-ca.crt --root-ca-file=/etc/kubernetes/pki/ca.crt --service-account-private-key-file=/etc/kubernetes/pki/sa.key --service-cluster-ip-range=10.96.0.0/12 --use-service-account-credentials=true
root    5053  0.1  1.4 1282960 56744 ?        Ssl  11:46   0:06 kube-scheduler --authentication-kubeconfig=/etc/kubernetes/scheduler.conf --authorization-kubeconfig=/etc/kubernetes/scheduler.conf --bind-address=127.0.0.1 --kubernetes/pki/sa.key --leader-elect=true
root    5081  1.4  7.4 1553796 287648 ?        Ssl  11:46   1:10 kube-apiserver --advertise-address=172.24.153.39 --allow-privileged=true --authorization-mode=Node,RBAC --client-ca-file=/etc/kubernetes/pki/ca.crt --enable-admission-plugins=NodeRestriction --enable-bootstrap-token-auth=true --etcd-cafile=/etc/kubernetes/pki/etcd/ca.crt --etcd-certfile=/etc/kubernetes/pki/apiserver-etcd-client.crt --etcd-keyfile=/etc/kubernetes/pki/apiserver-etcd-client.key --etcd-servers=https://127.0.0.1:2379 --kubelet-client-certificate=/etc/kubernetes/pki/apiserver-kubelet-client.crt --kubelet-client-key=/etc/kubernetes/pki/apiserver-kubelet-client.key --kubelet-preferred-address-types=InternalIP,ExternalIP,Hostname --proxy-client-cert-file=/etc/kubernetes/pki/front-proxy-client.crt --proxy-client-key-file=/etc/kubernetes/pki/front-proxy-client.key --requestheader-allowed-names=front-proxy-client --requestheader-client-ca-file=/etc/kubernetes/pki/front-proxy-ca.crt --requestheader-extra-headers-prefix=X-Remote-Extra- --requestheader-group-headers=X-Remote-Group --requestheader-username-headers=X-Remote-User --secure-port=6443 --service-account-issuer=https://kubernetes.default.svc.cluster.local --service-account-key-file=/etc/kubernetes/pki/sa.pub --service-account-signing-key-file=/etc/kubernetes/pki/sa.key --service-cluster-ip-range=10.96.0.0/12 --tls-cert-file=/etc/kubernetes/pki/apiserver.crt --tls-private-key-file=/etc/kubernetes/pki/apiserver.key
root    5089  0.6  1.4 11222012 57024 ?        Ssl  11:46   0:30 etcd --advertise-client-urls=https://172.24.153.39:2379 --cert-file=/etc/kubernetes/pki/etcd/server.crt --client-cert-auth=true --data-dir=/var/lib/etcd --experimental-initial-corrupt-check=true --experimental-watch-progress-notify-interval=5s --initial-advertise-peer-urls=https://172.24.153.39:2380 --initial-cluster=k8s-master=https://172.24.153.39:2380 --key-file=/etc/kubernetes/pki/etcd/server.key --listen-client-urls=https://127.0.0.1:2379,https://172.24.153.39:2379 --listen-metrics-urls=https://127.0.0.1:2381 --listen-peer-urls=https://172.24.153.39:2380 --name=k8s-master --peer-cert-file=/etc/kubernetes/pki/etcd/peer.crt --peer-client-cert-auth=true --peer-key-file=/etc/kubernetes/pki/etcd/peer.key --peer-trusted-ca-file=/etc/kubernetes/pki/etcd/ca.crt --snapshot-count=10000 --trusted-ca-file=/etc/kubernetes/pki/etcd/ca.crt
root    5171  0.7  2.6 2077828 103648 ?        Ssl  11:46   0:37 /usr/bin/kubelet --bootstrap-kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --kubeconfig=/etc/kubernetes/kubelet.conf --config=/var/lib/kubelet/config.yaml --container-runtime-endpoint=unix:///var/run/containerd/containerd.sock --pod-infra-container-image=registry.aliyuncs.com/google_containers/pause:3.9
root    5292  0.0  1.2 1281792 47428 ?        Ssl  11:46   0:00 /usr/local/bin/kube-proxy --config=/var/lib/kube-proxy/config.conf --hostname-override=k8s-master
root    6111  0.0  0.9 1269140 36728 ?        Ssl  12:23   0:01 /opt/bin/flanneld -ip-masq --kube-subnet-mgr
root    17637 0.0  0.0 9032 720 pts/0    S+   13:08   0:00 grep --color=auto kube
```

在worker节点上输出如下：

```
root@k8s-worker:/etc/containerd# ps aux | grep kube
root    4955  0.3  2.4 2077316 95332 ?        Ssl  12:29   0:08 /usr/bin/kubelet --bootstrap-kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --kubeconfig=/etc/kubernetes/kubelet.conf --config=/var/lib/kubelet/config.yaml --container-runtime-endpoint=unix:///var/run/containerd/containerd.sock --pod-infra-container-image=registry.aliyuncs.com/google_containers/pause:3.9
root    5104  0.0  1.2 1281792 48056 ?        Ssl  12:29   0:00 /usr/local/bin/kube-proxy --config=/var/lib/kube-proxy/config.conf --hostname-override=k8s-worker
root    5444  0.0  0.9 1267884 36460 ?        Ssl  12:31   0:01 /opt/bin/flanneld -ip-masq --kube-subnet-mgr
root    14679 0.0  0.0 9032 656 pts/0    S+   13:08   0:00 grep --color=auto kube
```

具体功能如下：

1.master节点上的 kube-controller-manager：Kube Controller Manager是一个守护进程，内嵌随 Kubernetes一起发布的核心控制回路。Kube Controller Manager 通过API服务器监控集群的状态，确保集群处于预期的工作状态。Kube Controller Manager 由负责不同资源的多个控制器构成。目前，Kubernetes自带的控制器包括副本控制器、节点控制器、命名空间控制器和服务账号控制器等。

(阿里云上的解释)

2.master节点上的 `kube-scheduler`：Kubernetes 调度器是一个控制面进程，负责将 Pods 指派到节点上。调度器基于约束和可用资源为调度队列中每个 Pod 确定其可合法放置的节点。调度器之后对所有合法的节点进行排序，将 Pod 绑定到一个合适的节点。

3.master节点上的 `kube-apiserver`：这个就是我们课上所解释的 `api-server`，API 服务器验证并配置 API 对象的数据，这些对象包括 pods、services、replicationcontrollers 等。API 服务器为 REST 操作提供服务，并为集群的共享状态提供前端，所有其他组件都通过该前端进行交互。

4.master节点上的 `etcd` 进程：`etcd` 是一个分布式键值存储系统，通常用于存储 Kubernetes 集群的配置信息、状态信息等数据。

5. `kubelet`：`kubelet` 是在每个 Node 节点上运行的主要“节点代理”。它可以使用以下之一向 `apiserver` 注册：主机名 (hostname)；覆盖主机名的参数；某云驱动的特定逻辑。

6. `kube-proxy`：这是k8s的网络代理，在每个节点上运行。网络代理反映了每个节点上 Kubernetes API 中定义的服务，并且可以执行简单的 TCP、UDP 和 SCTP 流转发，或者在一组后端进行循环 TCP、UDP 和 SCTP 转发。

7.flannel进程：`flannel` 是 Kubernetes 集群中一个网络管理组件，主要负责网络的管理和配置，与我们下载的flannel插件有关。

**Q3: 在两个节点中分别使用 `crictl ps` 显示所有正常运行的containerd容器，记录其输出，并简要说明各个容器所包含的k8s组件，以及那些k8s组件未运行在容器中**

回答：master节点的显示：

CONTAINER	IMAGE	CREATED	STATE	NAME	ATTEMPT	POD ID	POD
bb38dc090ff7	ead0a4a53df89	About an hour ago	Running	coredns	0	0d42c195080c3	coredns-66f779496c-mt8xm
92a1f939d1ca	ead0a4a53df89	About an hour ago	Running	coredns	0	c526a8d958e08	coredns-66f779496c-49ndq
b0c25e599ec3d	1575deaad3b05	About an hour ago	Running	kube-flannel	0	3a4c0847b3ff6	kube-flannel-ds-xwzsh
89be855de15fc	09c5e1abe5922	2 hours ago	Running	kube-proxy	0	3923b51881e2f	kube-proxy-nnc29
7327bf503af9a	3861cfd7c04c	2 hours ago	Running	etcd	0	f3bf8f385410b	etcd-k8s-master
a41cd6bcfec0	69947457eaa42	2 hours ago	Running	kube-apiserver	0	14c897041ede8	kube-apiserver-k8s-master
fb7e4a3dfc89f	7264907bfc5be	2 hours ago	Running	kube-scheduler	0	c452b5c0774e3	kube-scheduler-k8s-master
4508a142de6d2	8981bdcce6670	2 hours ago	Running	kube-controller-manager	0	706e2b1e7f946	kube-controller-manager-k8s-master

worker节点的显示：

CONTAINER	IMAGE	CREATED	STATE	NAME	ATTEMPT	POD ID	POD
bb972fdbafef5	1575deaad3b05	53 minutes ago	Running	kube-flannel	0	afaf811d1728a	kube-flannel-ds-22x65
60168d1406136	09c5e1abe5922	55 minutes ago	Running	kube-proxy	0	c541c7cf92aa1	kube-proxy-ctn44

`coredns` 组件：`CoreDNS` 是 Kubernetes 集群中一个常用的 DNS 服务器组件，用于解析主机名到 IP 地址。它主要负责为集群内部服务提供 DNS 解析服务，以便实现服务之间的互相发现和通信。

`kube-flannel` 组件：`kube-flannel` 是 Kubernetes 集群中一个常用的网络插件组件，用于实现网络的配置和管理。它通常与 `flannel`、`etcd` 等其他组件配合使用，以确保集群内网络的正确运行和通信。

其它组件的说明与Q2中的类似，这里不重复赘述。

像 `kubelet` 这种就没有运行在容器中，通常是直接运行在 Kubernetes 集群的节点上，它会监听 API 服务器的指令，执行 Pod 的创建、销毁、重启等操作，并监控容器的健康状态等。

**Q4: 请采用声明式接口对Pod进行部署，并将部署所需的yaml文件记录在实践文档中**

我们的test\_k8s.yaml文件如下：

```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: testpod
```



```

5     labels:
6       app: testpod
7   spec:
8     containers:
9       - name: fileserver
10        image: 7143192/fileserver:latest
11        ports:
12          - containerPort: 8080
13        volumeMounts:
14          - name: test-volume
15            mountPath: /usr/share/files
16       - name: downloader
17        image: 7143192/downloader:latest
18        ports:
19          - containerPort: 3000
20        volumeMounts:
21          - name: test-volume
22            mountPath: /data
23   volumes:
24     - name: test-volume
25       hostPath:
26         path: /home/data
27         type: DirectoryOrCreate

```

我们调用指令

```
1 kubectl apply -f test_k8s.yaml
```

一段时间后查看pod状态

```

root@k8s-master:/home# kubectl get pod
NAME          READY   STATUS    RESTARTS   AGE
testpod       2/2     Running   0           4m26s

```

可以看到有两个pod正在运行。

**Q5: 请在worker节点上，在部署Pod的前后分别采用 `crictl ps` 查看所有运行中的容器并对比两者的区别。请将创建该Pod所创建的全部新容器列举出来，并一一解释其作用**

回答：master节点上

CONTAINER	IMAGE	CREATED	STATE	NAME	ATTEMPT	POD ID	POD
3b638dc090ff7	ead0a4a53df89	2 hours ago	Running	coredns	0	0d42c195090c3	coredns-66f779496c-mt8xm
3261f0389d1ca	ead0a4a53df89	2 hours ago	Running	coredns	0	c526a8d955e88	coredns-66f779496c-49ndq
2bc25e5e9ec3d	1575deaad3b05	2 hours ago	Running	kube-flannel	0	3adc0047b3ff6	kube-flannel-ds-xwzsh
89be855de15fc	09c5e1abe5922	2 hours ago	Running	kube-proxy	0	3923b51881e2f	kube-proxy-nnc29
7327b5f03af9a	3861cfd7c04c	2 hours ago	Running	etcd	0	f3bf8f385410b	etcd-k8s-master
41cd60cfefc0	69947457aa42	2 hours ago	Running	kube-apiserver	0	14c8970a1ade8	kube-apiserver-k8s-master
fb7e4a3dfc09f	7264907bfc5be	2 hours ago	Running	kube-scheduler	0	c452b5c0774e3	kube-scheduler-k8s-master
4505a142de6d2	8981bddce6670	2 hours ago	Running	kube-controller-manager	0	706e2b1e7f946	kube-controller-manager-k8s-master

worker节点上

CONTAINER	IMAGE	CREATED	STATE	NAME	ATTEMPT	POD ID	POD
b971dd036333d	eea933b9ed0ff	6 minutes ago	Running	downloader	0	e6cf53e8a57e5	testpod
2bb4d8a64ff73	e6a8b3fdc6504	6 minutes ago	Running	fileserver	0	e6cf53e8a57e5	testpod
bb972fdbafef5	1575deaad3b05	About an hour ago	Running	kube-flannel	0	afaf811d1728a	kube-flannel-ds-22x65
60168d1406136	09c5e1abe5922	2 hours ago	Running	kube-proxy	0	c541c7cf92aa1	kube-proxy-ctn44

与我们在Q3中打印出的信息对比，发现master节点上没有变化，新创建的pod运行在了worker节点上。容器的作用在前面已经解释过，这里只解释worker节点新增的两个节点的作用：

如这两个容器的名字所展示的一样，一个是downloader，另外一个为fileserver。



这里很奇怪的一点是没有pause容器，Pause容器是Kubernetes系统自动生成的一个特殊容器，通常用于实现Pod的网络命名空间共享和挂载共享卷，在Kubernetes中起到连接和协调其他容器的作用，确保Pod能够正常运行并实现容器之间的共享和通信。

经过网上的查阅，在<https://www.saonihuo.com/question/detail-2150987.html>中查到了原因，

我们用 `ctr -n k8s.io c ls` 查看所有运行在k8s上的 containerd 容器，显示如下：

```
root@k8s-worker:/home# ctr -n k8s.io c ls
WARN[0000] DEPRECATION: CRI API v1alpha2 is deprecated since containerd v1.7 and removed in containerd v2.0. Use CRI API v1 instead.
CONTAINER                                IMAGE                                RUNTIME
60168d14061363098b4d43c6c259675fa77a6623719c20913b6831572976aad registry.aliyuncs.com/google_containers/kube-proxy:v1.28.9 io.containerd.runc.v2
7740df2af8533d1ef4632ebf5d7821217400b53542f5b1de9210c20f090847 docker.io/7143192/downloader:latest io.containerd.runc.v2
81a0dd5b39423def77cad132cffe02a40887572b8b270f965c42b3c38c00a0d18 docker.io/7143192/fileserver:latest io.containerd.runc.v2
afaf811d1728ab4517bf091b6b6c47e81fc6465826550cb542bc7855746fa01b registry.aliyuncs.com/google_containers/pause:3.9 io.containerd.runc.v2
bb3f4630a3472c9472a7e2ffa5f9d92fa85f7c3aa833b45530b2941d06d790f6 registry.aliyuncs.com/google_containers/pause:3.9 io.containerd.runc.v2
bb972fdbafef5f21077bba0076d2cc4697097365e3be1d49e9a092140fab0ed4 docker.io/flannel/flannel:v0.25.1 io.containerd.runc.v2
c541c7c792aa18e35c862c0e640cee429d82b862b5b8aeca23037f0100ca465 registry.aliyuncs.com/google_containers/pause:3.9 io.containerd.runc.v2
e0ea3220788927eee5a969efc40dd0e3ab0954b63d5022775e12fdd50fab8a7e docker.io/flannel/flannel-cni-plugin:v1.4.1-flannel io.containerd.runc.v2
e0ab55e0c1cc306dfc0d2d3e5ef816e1359ca2d4459e032bad27aa9cd4214a registry.aliyuncs.com/google_containers/pause:3.9 io.containerd.runc.v2
f76959b2095b80b3c0027525ad68b3e9a5b58efbe6a0223427556346c968f8c docker.io/flannel/flannel:v0.25.1 io.containerd.runc.v2
```

这样就看到了我们的pause容器，所以pause容器也是存在的。

**Q6: 请结合博客 [https://blog.51cto.com/u\\_15069443/4043930](https://blog.51cto.com/u_15069443/4043930) 的内容，将容器中的veth与host机器上的veth匹配起来，并采用 ip link 和 ip addr 指令找到位于host机器中的所有网络设备及其之间的关系。结合两者的输出，试绘制出worker节点中涉及新部署Pod的所有网络设备和其网络结构，并在图中标注出从master节点中使用pod ip访问位于worker节点中的Pod的网络路径**

回答：通过以下指令进入容器

```
1 # 这个是downloader容器
2 crictl exec -it b971dd036333d /bin/bash
3 # 这个是fileservers容器
4 crictl exec -it 2bb4d8a64ff73 /bin/bash
```

进入后执行

```
1 apt-get update
2 apt-get install -y iproute2
```

安装对应的工具，然后参照博客

downloader中：

```
root@testpod:/apps# ip link show eth0
2: eth0@if5: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1450 qdisc noqueue state UP mode DEFAULT group default
    link/ether be:ab:81:b0:b7:72 brd ff:ff:ff:ff:ff:ff link-netnsid 0
root@testpod:/apps#

root@testpod:/apps# ip addr
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: eth0@if5: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1450 qdisc noqueue state UP group default
    link/ether be:ab:81:b0:b7:72 brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 10.244.1.2/24 brd 10.244.1.255 scope global eth0
        valid_lft forever preferred_lft forever
    inet6 fe80::bcab:81ff:feb0:b772/64 scope link
        valid_lft forever preferred_lft forever
```

fileservers中：

```
root@testpod:/apps# ip link show eth0
2: eth0@if5: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1450 qdisc noqueue state UP mode DEFAULT group default
    link/ether be:ab:81:b0:b7:72 brd ff:ff:ff:ff:ff:ff link-netnsid 0
root@testpod:/apps#
```

```

root@testpod:/apps# ip addr
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: eth0@if5: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1450 qdisc noqueue state UP group default
    link/ether be:ab:81:b0:b7:72 brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 10.244.1.2/24 brd 10.244.1.255 scope global eth0
        valid_lft forever preferred_lft forever
    inet6 fe80::bcab:81ff:feb0:b772/64 scope link
        valid_lft forever preferred_lft forever
root@testpod:/apps#

```

在host主机中:

```

root@k8s-worker:/etc/containerd# ip link show | grep 5
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP mode DEFAULT group default qlen 1000
3: flannel.1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1450 qdisc noqueue state UNKNOWN mode DEFAULT group default
    link/ether 36:ac:8f:79:59:c0 brd ff:ff:ff:ff:ff:ff
4: cni0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1450 qdisc noqueue state UP mode DEFAULT group default qlen 1000
    link/ether 1a:11:75:60:fd:05 brd ff:ff:ff:ff:ff:ff
5: veth05dbc1b4@if2: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1450 qdisc noqueue master cni0 state UP mode DEFAULT group default
    link/ether 02:2a:2d:32:52:9a brd ff:ff:ff:ff:ff:ff link-netns cni-a4de7ab4-a30f-f2fc-e527-9f01229125dc
root@k8s-worker:/etc/containerd#

```

```

root@k8s-worker:/etc/containerd# ip addr
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP group default qlen 1000
    link/ether 00:16:3e:26:d1:f6 brd ff:ff:ff:ff:ff:ff
    inet 172.23.193.146/20 brd 172.23.207.255 scope global dynamic eth0
        valid_lft 315340506sec preferred_lft 315340506sec
    inet6 fe80::216:3eff:fe26:d1f6/64 scope link
        valid_lft forever preferred_lft forever
3: flannel.1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1450 qdisc noqueue state UNKNOWN group default
    link/ether 36:ac:8f:79:59:c0 brd ff:ff:ff:ff:ff:ff
    inet 10.244.1.0/32 scope global flannel.1
        valid_lft forever preferred_lft forever
    inet6 fe80::34ac:8fff:fe79:59c0/64 scope link
        valid_lft forever preferred_lft forever
4: cni0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1450 qdisc noqueue state UP group default qlen 1000
    link/ether 1a:11:75:60:fd:05 brd ff:ff:ff:ff:ff:ff
    inet 10.244.1.1/24 brd 10.244.1.255 scope global cni0
        valid_lft forever preferred_lft forever
    inet6 fe80::1811:75ff:fe60:fd05/64 scope link
        valid_lft forever preferred_lft forever
5: veth05dbc1b4@if2: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1450 qdisc noqueue master cni0 state UP group default
    link/ether 02:2a:2d:32:52:9a brd ff:ff:ff:ff:ff:ff link-netns cni-a4de7ab4-a30f-f2fc-e527-9f01229125dc
    inet6 fe80::2a:2dff:fe32:529a/64 scope link
        valid_lft forever preferred_lft forever

```

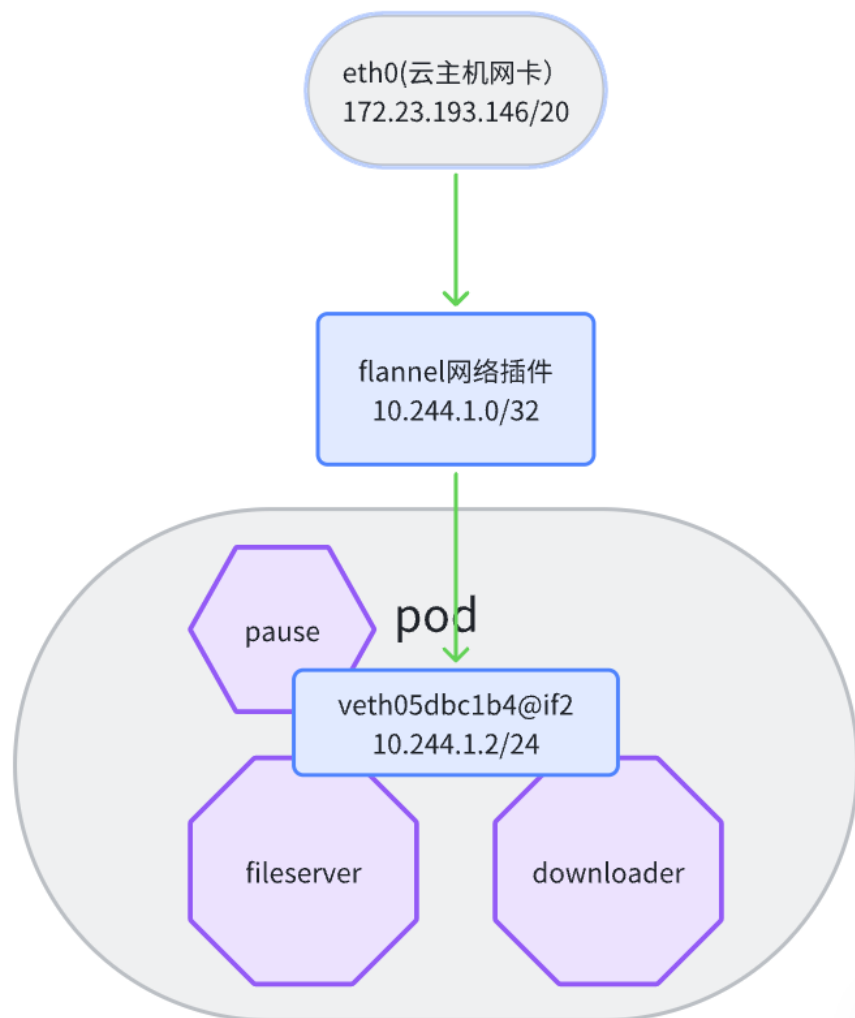
查看对应5的 vethinterface 是哪一个

好, 这里我们就知道这是 veth-pair 技术, 是一对虚拟的网络设备接口, 成对出现。一端链接主机, 一端链接容器, 这样可以让它们之间直接通讯。

我们在上面主机执行的时候, 看到的是 5: veth05dbc1b4@if2, 2是它pair配对的 veth 的索引index, 这里就是我们之前在两个容器之中看到的 2: eth0@if5, 刚好我们发现它配对的veth的索引是5, 与上面的恰好顺序相反。

同一个网段间的设备之间可以相互通信, 我们在主机中可以很轻松找到flannel网络插件对应的虚拟网卡, 另外, 流量的转发也一定离不开服务器自带的虚拟网卡eth0(节点的网络接口), 对应服务器的内网 ip

拓扑图大致如下:



**Q7: 请采用声明式接口对Deployment进行部署，并将Deployment所需要的yaml文件记录在文档中**

回答：我们的 `test_k8s_deployment.yaml` 如下写法：

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: test-deployment
5  spec:
6    replicas: 3
7    selector:
8      matchLabels:
9        app: testpod
10   template:
11     metadata:
12       labels:
13         app: testpod
14     spec:
15       containers:
16         - name: fileserver
17           image: 7143192/fileserver:latest
18           ports:
19             - containerPort: 8080
20       volumeMounts:
21         - name: test-volume
22           mountPath: /usr/share/files
```

```

23     - name: downloader
24       image: 7143192/downloader:latest
25       ports:
26         - containerPort: 3000
27       volumeMounts:
28         - name: test-volume
29           mountPath: /data
30     volumes:
31     - name: test-volume
32       hostPath:
33         path: /home/data2
34         type: DirectoryOrCreate

```

我们在master节点上使用如下指令：

```

1 root@k8s-master:/home# kubectl apply -f test_k8s_deployment.yaml
2 deployment.apps/test-deployment created
3 root@k8s-master:/home# kubectl get deployment
4 NAME                                READY    UP-TO-DATE    AVAILABLE    AGE
5 test-deployment                    3/3      3              3            23s

```

我们可以看到成功创建

通过get pod

```

1 root@k8s-master:/home# kubectl get pods
2 NAME                                READY    STATUS    RESTARTS    AGE
3 test-deployment-754f9f879-pzlv7     2/2      Running   0            34s
4 test-deployment-754f9f879-rfm6b     2/2      Running   0            34s
5 test-deployment-754f9f879-tprpj     2/2      Running   0            34s
6 testpod                             2/2      Running   0            94m

```

因为我们设置的replicas为3，所以会产生3个pod。

**Q8: 在该使用Deployment的部署方式下，不同Pod之间的文件是否共享？该情况会在实际使用文件下载与共享服务时产生怎样的实际效果与问题？应如何解决这一问题？**

回答：默认情况下容器的文件系统是互相隔离的，但是我们在yaml文件中声明了需要绑定的Volume，并且我们将其挂载在了主机下，所以最后会导致挂载的那一部分目录内容是共享的。

这可能会导致冲突产生，如文件名冲突（你写的一个文件名可能已经被别人写了，那么就不能以这个文件名在作为文件了），甚至一个用户（对于某个特定的downloader）可以下载到别人的文件（可能是别的fileserver写的），因为数据共享，导致安全性问题。所以我们需要使用PV（持久化卷）解决上面的问题，它是对底层共享存储的一种抽象，将共享存储定义为一种资源，不属于任何的namespace，属于集群级别资源。用户使用PV需要通过PersistentVolumeClaim（PVC），我们在使用PV可以通过设置不同的参数来解决共享同一个volume的问题。

例如下面这样：

```

1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: unique-volume-deployment
5 spec:
6   replicas: 2

```

```

7   selector:
8     matchLabels:
9       app: unique-volume
10  template:
11    metadata:
12      labels:
13        app: unique-volume
14    spec:
15      containers:
16        - name: unique-volume-container
17          image: your-image
18          volumeMounts:
19            - name: unique-volume
20              mountPath: /path/to/data
21      volumes:
22        - name: unique-volume
23          persistentVolumeClaim:
24            claimName: unique-pvc-$(podname) # 使用 podname 动态生成不同的 PVC
名称

```

```

1  apiVersion: v1
2  kind: PersistentVolumeClaim
3  metadata:
4    name: unique-pvc-$(podname) # 使用 podname 动态生成不同的 PVC 名称
5  spec:
6    accessModes:
7      - ReadWriteOnce
8    resources:
9      requests:
10       storage: 1Gi

```

这样就能根据不同的podname生成不同的PVC从而每个pod共享自己独有的volume

### Q9: 请采用声明式接口对Service进行部署，并将部署所需的yaml文件记录在实践文档中

回答：我们的testservice.yaml文件如下：

```

1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: test-service
5  spec:
6    selector:
7      app: testpod
8    ports:
9      - port: 8080
10       targetPort: 8080
11       name: fileserver
12      - port: 3000
13       targetPort: 3000
14       name: downloader

```

在master节点上运行如下指令：

```
root@k8s-master:/home# kubectl apply -f test-service.yaml
service/test-service created
root@k8s-master:/home# kubectl get service
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	4h4m
test-service	ClusterIP	10.101.163.22	<none>	8080/TCP,3000/TCP	8s

我们完成了名为test-service服务的部署。

通过

```
root@k8s-master:/home# kubectl get endpoints test-service -o wide
```

NAME	ENDPOINTS	AGE
test-service	10.244.1.3:8080,10.244.1.4:8080,10.244.1.5:8080 + 3 more...	3m48s

也可以看到绑定对应的endpoints成功。

**Q10: 请在master节点中使用 iptables-save 指令输出所有的iptables规则，将其中与Service访问 相关的iptables规则记录在实践文档中，并解释网络流量是如何采用基于iptables的方式被从对Service 的clusterIP的访问定向到实际的Pod中的，又是如何实现负载均衡到三个pod的。**

回答：调用指令终端打印如下：

```
root@k8s-master:/home# sudo iptables-save | grep test-service
-A KUBE-SEP-2H3DDHMGCONVHAPW -s 10.244.1.3/32 -m comment --comment "default/test-service:downloader" -j KUBE-MARK-MASQ
-A KUBE-SEP-2H3DDHMGCONVHAPW -p tcp -m comment --comment "default/test-service:downloader" -m tcp -j DNAT --to-destination 10.244.1.3:3000
-A KUBE-SEP-B2WJKGR0Q4DNXQXB -s 10.244.1.3/32 -m comment --comment "default/test-service:fileserver" -j KUBE-MARK-MASQ
-A KUBE-SEP-B2WJKGR0Q4DNXQXB -p tcp -m comment --comment "default/test-service:fileserver" -m tcp -j DNAT --to-destination 10.244.1.3:8080
-A KUBE-SEP-QJP363NGI67SEY4S -s 10.244.1.4/32 -m comment --comment "default/test-service:downloader" -j KUBE-MARK-MASQ
-A KUBE-SEP-QJP363NGI67SEY4S -p tcp -m comment --comment "default/test-service:downloader" -m tcp -j DNAT --to-destination 10.244.1.4:3000
-A KUBE-SEP-RBA7ORS5GEW307HL4 -s 10.244.1.4/32 -m comment --comment "default/test-service:fileserver" -j KUBE-MARK-MASQ
-A KUBE-SEP-RBA7ORS5GEW307HL4 -p tcp -m comment --comment "default/test-service:fileserver" -m tcp -j DNAT --to-destination 10.244.1.4:8080
-A KUBE-SEP-SCOR00BP12P6R6Y -s 10.244.1.5/32 -m comment --comment "default/test-service:fileserver" -j KUBE-MARK-MASQ
-A KUBE-SEP-SCOR00BP12P6R6Y -p tcp -m comment --comment "default/test-service:fileserver" -m tcp -j DNAT --to-destination 10.244.1.5:8080
-A KUBE-SEP-X7S6D3RE5G7MRR2 -s 10.244.1.5/32 -m comment --comment "default/test-service:downloader" -j KUBE-MARK-MASQ
-A KUBE-SEP-X7S6D3RE5G7MRR2 -p tcp -m comment --comment "default/test-service:downloader" -m tcp -j DNAT --to-destination 10.244.1.5:3000
-A KUBE-SERVICES -d 10.101.163.22/32 -p tcp -m comment --comment "default/test-service:fileserver cluster IP" -m tcp --dport 8080 -j KUBE-SVC-YPW4ESSQ7BDH6RZX
-A KUBE-SERVICES -d 10.101.163.22/32 -p tcp -m comment --comment "default/test-service:downloader cluster IP" -m tcp --dport 3000 -j KUBE-SVC-5DBF63XQ64WV75V6
-A KUBE-SVC-5DBF63XQ64WV75V6 -s 10.244.0.0/16 -d 10.101.163.22/32 -p tcp -m comment --comment "default/test-service:downloader cluster IP" -m tcp --dport 3000 -j KUBE-MARK-MASQ
-A KUBE-SVC-5DBF63XQ64WV75V6 -m comment --comment "default/test-service:downloader" --> 10.244.1.3:3000 -m statistic --mode random --probability 0.3333333349 -j KUBE-SEP-2H3DDHMGCONVHAPW
-A KUBE-SVC-5DBF63XQ64WV75V6 -m comment --comment "default/test-service:downloader" --> 10.244.1.4:3000 -m statistic --mode random --probability 0.5000000000 -j KUBE-SEP-QJP363NGI67SEY4S
-A KUBE-SVC-5DBF63XQ64WV75V6 -m comment --comment "default/test-service:downloader" --> 10.244.1.5:3000 -j KUBE-SEP-X7S6D3RE5G7MRR2
-A KUBE-SVC-YPW4ESSQ7BDH6RZX -s 10.244.0.0/16 -d 10.101.163.22/32 -p tcp -m comment --comment "default/test-service:fileserver cluster IP" -m tcp --dport 8080 -j KUBE-MARK-MASQ
-A KUBE-SVC-YPW4ESSQ7BDH6RZX -m comment --comment "default/test-service:fileserver" --> 10.244.1.3:8080 -m statistic --mode random --probability 0.3333333349 -j KUBE-SEP-B2WJKGR0Q4DNXQXB
-A KUBE-SVC-YPW4ESSQ7BDH6RZX -m comment --comment "default/test-service:fileserver" --> 10.244.1.4:8080 -m statistic --mode random --probability 0.5000000000 -j KUBE-SEP-RBA7ORS5GEW307HL4
-A KUBE-SVC-YPW4ESSQ7BDH6RZX -m comment --comment "default/test-service:fileserver" --> 10.244.1.5:8080 -j KUBE-SEP-SCOR00BP12P6R6Y
```

这里我们可以看到有三个比较关键的部分： KUBE\_SVC , KUBE\_SEP , KUBE\_SERVICES

首先 KUBE\_SERVICES 是位于最外层，暴露服务的端口，两行规则如下：

```
1 -A KUBE-SERVICES -d 10.101.163.22/32 -p tcp -m comment --comment
  "default/test-service:fileserver cluster IP" -m tcp --dport 8080 -j KUBE-SVC-
  YPW4ESSQ7BDH6RZX
2 -A KUBE-SERVICES -d 10.101.163.22/32 -p tcp -m comment --comment
  "default/test-service:downloader cluster IP" -m tcp --dport 3000 -j KUBE-SVC-
  5DBF63XQ64WV75V6
```

第一行规则指定了当目标IP地址为10.101.163.22且目标TCP端口为8080时，应该将流量重定向到名为 KUBE-SVC-YPW4ESSQ7BDH6RZX的规则进行处理。这个规则对应着test-service的fileserver服务。第二行规则则指定了当目标IP地址为10.101.163.22且目标TCP端口为3000时，应该将流量重定向到名为 KUBE-SVC-5DBF63XQ64WV75V6的规则进行处理。这个规则对应着test-service的downloader服务。

其次是 KUBE\_SVC ，它的八行规则如下：



```

1 -A KUBE-SVC-5DBF63XQ64WV75V6 ! -s 10.244.0.0/16 -d 10.101.163.22/32 -p tcp -m comment --comment "default/test-service:downloader cluster IP" -m tcp --dport 3000 -j KUBE-MARK-MASQ
2 -A KUBE-SVC-5DBF63XQ64WV75V6 -m comment --comment "default/test-service:downloader -> 10.244.1.3:3000" -m statistic --mode random --probability 0.33333333349 -j KUBE-SEP-2H3DDHW6CONVHAPW
3 -A KUBE-SVC-5DBF63XQ64WV75V6 -m comment --comment "default/test-service:downloader -> 10.244.1.4:3000" -m statistic --mode random --probability 0.50000000000 -j KUBE-SEP-OJP363NGI67SEY4S
4 -A KUBE-SVC-5DBF63XQ64WV75V6 -m comment --comment "default/test-service:downloader -> 10.244.1.5:3000" -j KUBE-SEP-X7S6D3RE5GM7MRR2
5 -A KUBE-SVC-YPW4ESSQ7BDH6RZX ! -s 10.244.0.0/16 -d 10.101.163.22/32 -p tcp -m comment --comment "default/test-service:fileserver cluster IP" -m tcp --dport 8080 -j KUBE-MARK-MASQ
6 -A KUBE-SVC-YPW4ESSQ7BDH6RZX -m comment --comment "default/test-service:fileserver -> 10.244.1.3:8080" -m statistic --mode random --probability 0.33333333349 -j KUBE-SEP-B2WJK6ROO4DNXQXB
7 -A KUBE-SVC-YPW4ESSQ7BDH6RZX -m comment --comment "default/test-service:fileserver -> 10.244.1.4:8080" -m statistic --mode random --probability 0.50000000000 -j KUBE-SEP-RBA7ORSGEW3O7HL4
8 -A KUBE-SVC-YPW4ESSQ7BDH6RZX -m comment --comment "default/test-service:fileserver -> 10.244.1.5:8080" -j KUBE-SEP-SCOROOBPIJ2P6R6Y
9

```

这里上面四条是处理downloader的，后面四条是处理fileserver的。

我们可以看到的是，流量转发到不同pod上的概率，就以前四条处理downloader的来说明，当流量到来时，转移到10.244.1.3:3000的概率为0.333（由probability 0.33333333349可知，因为有三个节点，所以概率是1/3），剩下0.67的概率会继续执行下面的规则，接着是转发到10.244.1.4: 3000的概率为0.5（这里是我们的前一条规则以0.67的概率继续向下执行，这个时候有两个节点，所以概率是0.5），剩下0.5的概率继续执行下面的规则，最后是转发到10.244.1.5: 3000的概率，这个时候已经经过前面两个规则（已经不会转发到10.244.1.3:3000和10.244.1.4: 3000），剩下的概率自然是100%，上面的过程基本上保证了每个pod的负载均衡，基本是1/3，**这种均衡是随机的负载均衡，不是基于IP-Hash，也不是round-robin**。在决定好转发到那个pod后，可以看到这三行规则每个后面都有一个类似 -j KUBE-SEP-2H3DDHW6CONVHAPW，也就是根据具体的pod转发到对应的 KUBE-SEP，KUBE-SEP 是 KUBE-SVC 对应的终端，它后面跟着具体的DNAT规则，例如：

```

1 -A KUBE-SEP-2H3DDHW6CONVHAPW -p tcp -m comment --comment "default/test-service:downloader" -m tcp -j DNAT --to-destination 10.244.1.3:3000
2

```

当流量匹配到这条规则时，会触发对应的处理动作。处理动作是进行目的地址网络地址转换（DNAT），将流量重定向到目标地址为10.244.1.3，目标端口为3000的目的IP地址。

简单来说，就是下面这样：

1.每个Service的每个服务端口都会在Chain KUBE-SERVICES中有一条对应的规则，发送到 clusterIP 的报文，将会转发到对应的Service的规则链，没有命中 clusterIP 的，转发到KUBE-NODEPORTS。

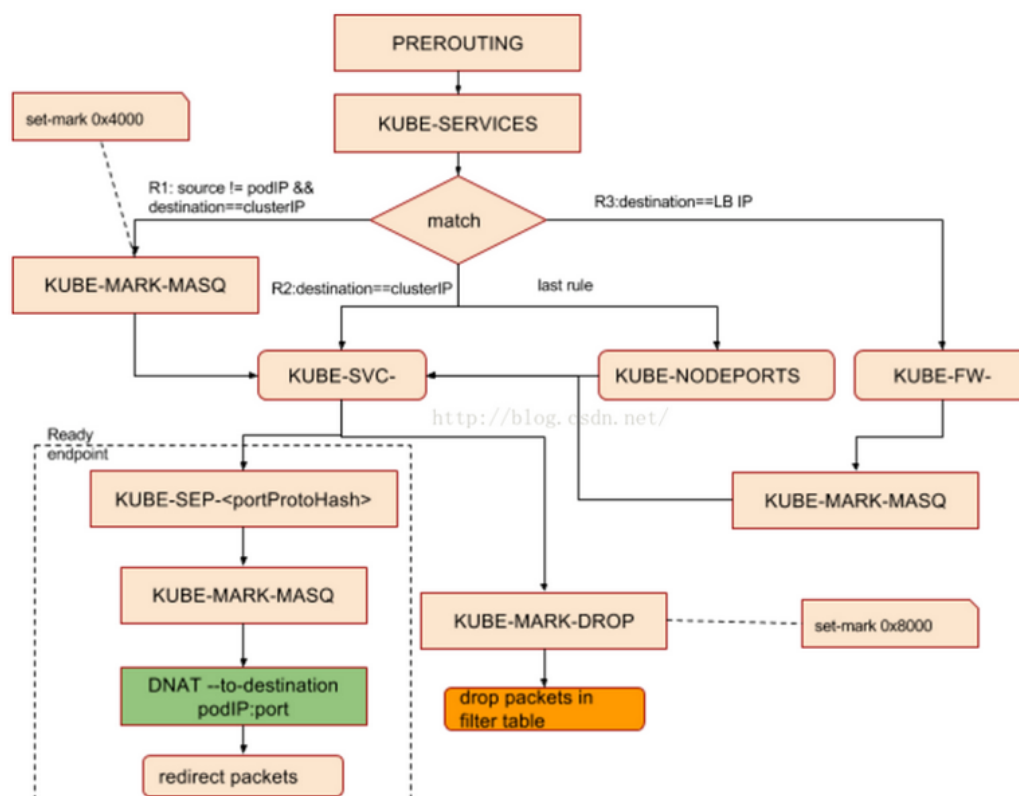
只有发送到被 kubernetes 占用的端口的报文才会进入KUBE-MARK-MASQ打上标记，并转发到对应的服务规则链。

2.每一个KUBE\_SERVICE，又将报文提交到了各自的KUBE-SEP-XXX



3.最后在KUBE-SEP-XX中完整了最终的DNAT，将目的地址转换成了POD的IP和端口。  
这里的KUBE-MARK-MASQ为报文打上了标记，表示这个报文是由 kubernetes 管理的，kubernetes 将会对它进行NAT转换。

流程图像这样：



### Q11: kube-proxy组件在整个service的定义与实现过程中起到了什么作用？请自行查找资料，并解释 在iptables模式下， kube-proxy的功能

回答： kube-proxy 会部署在k8s集群中的每个node节点上，它能够实现k8s service的通信和负载均衡， kube-proxy 负责帮助pod创建代理服务，从k8s的 api-server 获取server信息，并且根据信息创建代理服务，实现了从服务器端到pod的请求路由与转发，在iptables模式下， kube-proxy 的功能如下：

1. 实现Service的负载均衡： kube-proxy 会在节点上为每个Service创建相应的iptables规则，从而实现Service的负载均衡功能。这些规则会根据Service的类型（ClusterIP、NodePort、LoadBalancer）来进行配置，确保流量能够正确地路由到后端Pod上。
2. 实现Service的访问策略： kube-proxy 也会根据Service的定义，实现相应的访问策略，如 SessionAffinity 策略，确保请求会被定向到相同的后端Pod上。
3. 实现Service的可达性： kube-proxy 会监视后端Pod的健康状态，并根据实际情况更新相应的iptables规则，以确保流量不会被发送到不可达的后端Pod上。
4. 实现Service的服务发现： kube-proxy 会监视Service和Endpoints的变化，随时更新iptables规则以确保新的Pod能够被正确地访问到。
5. 定时从 etcd 服务获取到service信息来做相应的策略，维护网络规则和四层负载均衡工作
6. 如同我们上面的例子， kube-proxy 管理服务endpoints，并且让这个service对外暴露一个虚拟 ip，也就是 CLUSTER-IP，集群内部可以通过访问这个 ip 对应的端口就能顺利转发访问到集群内对应service的pod

**Q12: 请在上面部署的Deployment的基础上为其配置HPA，并将部署所需的yaml文件记录在实践文档中，如果对上面的Deployment配置有修改也请在文档中说明。具体参数为最大副本数为6，最小副本数为3，目标cpu平均利用率为40%。**

回答：首先我们在使用 hpa 之前需要先安装 Metrics Service插件。

安装过程如下：

```
1 root@k8s-master:/home# wget https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml
2
```

这里我们使用 wget 拉取 components.yaml 文件，我们需要对里面的一些内容做一些修改（最主要的是修改镜像源）。

修改这个文件的对应内容如下（只展示修改部分）：

```
1 spec:
2   containers:
3   - args:
4     - --cert-dir=/tmp
5     - --secure-port=10250
6     - --kubelet-preferred-address-types=InternalIP,ExternalIP,Hostname
7     - --kubelet-use-node-status-port
8     - --kubelet-insecure-tls
9     - --metric-resolution=15s
10    image: registry.cn-hangzhou.aliyuncs.com/google_containers/metrics-server:v0.7.1
```

修改源为阿里，同时需要加上 kubelet-insecure-tls 不验证客户端证书。

接着运行：

```
1 root@k8s-master:/home# kubectl apply -f components.yaml
2 serviceaccount/metrics-server created
3 clusterrole.rbac.authorization.k8s.io/system:aggregated-metrics-reader created
4 clusterrole.rbac.authorization.k8s.io/system:metrics-server created
5 rolebinding.rbac.authorization.k8s.io/metrics-server-auth-reader created
6 clusterrolebinding.rbac.authorization.k8s.io/metrics-server:system:auth-delegator created
7 clusterrolebinding.rbac.authorization.k8s.io/system:metrics-server created
8 service/metrics-server created
9 deployment.apps/metrics-server created
10 apiservice.apiregistration.k8s.io/v1beta1.metrics.k8s.io created
11
```

然后我们用 get pod 查看并且利用 top node 和 top pod 查看是否安装成功

```

root@k8s-master:/home# kubectl get pods -n kube-system
NAME                                READY    STATUS    RESTARTS   AGE
coredns-66f779496c-49ndq           1/1     Running   0           10h
coredns-66f779496c-mt8xm           1/1     Running   0           10h
etcd-k8s-master                     1/1     Running   0           10h
kube-apiserver-k8s-master           1/1     Running   0           10h
kube-controller-manager-k8s-master 1/1     Running   0           10h
kube-proxy-ctn44                    1/1     Running   0           9h
kube-proxy-nnc29                    1/1     Running   0           10h
kube-scheduler-k8s-master           1/1     Running   0           10h
metrics-server-997f546df-9s5mq      1/1     Running   0           36s
root@k8s-master:/home# kubectl top node
NAME      CPU(cores)   CPU%    MEMORY(bytes)   MEMORY%
k8s-master  52m          2%      1382Mi          37%
k8s-worker  19m          0%      943Mi           25%
root@k8s-master:/home# kubectl top pod
NAME                                CPU(cores)   MEMORY(bytes)
test-deployment-754f9f879-pzlv7     0m           3Mi
test-deployment-754f9f879-rfm6b     1m           3Mi
test-deployment-754f9f879-tpnpj     0m           3Mi
root@k8s-master:/home#

```

至此Metrics Service插件安装成功。

下面来执行本部分的任务，我们修改原来的 `test_k8s_deployment.yaml` 文件如下：

```

1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: test-deployment
5  spec:
6    # replicas: 3
7    selector:
8      matchLabels:
9        app: testpod
10   template:
11     metadata:
12       labels:
13         app: testpod
14     spec:
15       containers:
16         - name: fileserver
17           image: 7143192/fileserver:latest
18           ports:
19             - containerPort: 8080
20           volumeMounts:
21             - name: test-volume
22               mountPath: /usr/share/files
23           resources:
24             limits:
25               cpu: "500m"
26             requests:
27               cpu: "200m"
28         - name: downloader
29           image: 7143192/downloader:latest
30           ports:
31             - containerPort: 3000
32           volumeMounts:
33             - name: test-volume
34               mountPath: /data
35           resources:
36             limits:
37               cpu: "500m"
38             requests:
39               cpu: "200m"
40   volumes:

```

```

41     - name: test-volume
42       hostPath:
43         path: /home/data2
44         type: DirectoryOrCreate

```

本质上的修改是我去掉了replicas字段，同时给每个container的资源分配做了一个要求和限制。

下面是test\_hpa.yaml文件的内容：

```

1  apiVersion: autoscaling/v2
2  kind: HorizontalPodAutoscaler
3  metadata:
4    name: test-hpa
5  spec:
6    scaleTargetRef:
7      apiVersion: apps/v1
8      kind: Deployment
9      name: test-deployment
10   minReplicas: 3
11   maxReplicas: 6
12   metrics:
13     - type: Resource
14       resource:
15         name: cpu
16         target:
17           type: Utilization
18           averageUtilization: 40
19   behavior:
20     scaleDown:
21       policies:
22         - type: Percent
23           value: 10
24           periodSeconds: 60

```

这里我们已经对Q12和Q13两个问题的要求都已经写入yaml文件中，其中Q12的为具体参数为最大副本数为6，最小副本数为3，目标cpu平均利用率为40%，已经写在scaleTargetRef字段下，Q13的为hpa配置缩容的速率限制为每分钟10%，已经写在behavior字段下。

运行时重新apply之前deployment的yaml文件后运行

```

root@k8s-master:/home# kubectl apply -f test_hpa.yaml
horizontalpodautoscaler.autoscaling/test-hpa created

```

最后通过get hpa查看

```

root@k8s-master:/home# kubectl get hpa
NAME          REFERENCE                TARGETS  MINPODS  MAXPODS  REPLICAS  AGE
test-hpa      Deployment/test-deployment  0%/40%   3         6         3          68s

```

**Q13:小明发现，hpa缩容过快在负载抖动的情况下会引起pod被不必要地删除和再次创建，因此他决定限制缩容的速率，请为hpa配置缩容的速率限制为每分钟10%，并将部署所需的yaml文件记录在实践文档中**

回答：配置文件test\_hpa.yaml Q12已经给出，相应设置已经在其中。现在我们想要感受一下HPA带来的水平伸缩，所以我们需要给我们的集群上一些压力。

我们运行下面的指令：

```
1 kubectl run "pod-test0" --image=busybox:latest --restart=Never -- /bin/sh -c
  "while sleep 0.01; do wget -q -O- http://10.101.163.22:8080; done"
```

这个命令将创建一个临时的Pod对象，名称为"pod-test0"，使用 busybox:latest 镜像，并且设置不重新启动。这个Pod将会执行一个命令，不停地访问<http://10.101.163.22:8080>这个URL，可以测试<http://10.101.163.22:8080>上的负载。

```
root@k8s-master:/home# kubectl get pod
NAME                                READY    STATUS    RESTARTS    AGE
pod-test0                          1/1      Running   0            18s
test-deployment-5c65fd9775-h7grg    2/2      Running   4 (80s ago)  17m
test-deployment-5c65fd9775-s2lxb    2/2      Running   4 (87s ago)  28m
test-deployment-5c65fd9775-tvvfm    2/2      Running   4 (82s ago)  17m
root@k8s-master:/home#
```

可以看到的时已经在跑了，

```
root@k8s-master:/home# kubectl get hpa
NAME      REFERENCE                TARGETS  MINPODS  MAXPODS  REPLICAS  AGE
test-hpa  Deployment/test-deployment  1%/40%   3         6         3          18m
```

我们发现CPU利用率有所变化，但是显然还不够明显。

所以我们用下面的脚本test\_add.sh：

```
1 for i in {1..80}
2 do
3   podname="test$i"
4   kubectl run "pod-$podname" --image=busybox:latest --restart=Never --
  /bin/sh -c "while sleep 0.01;do wget -q -O- http://10.101.163.22:8080; done"
5 done
```

逻辑也非常简单，就是创建80个节点不停地访问<http://10.101.163.22:8080>这个URL，增加上面的负载。

同时我们修改 test\_hpa.yaml ,修改如下：

```
1 apiVersion: autoscaling/v2
2 kind: HorizontalPodAutoscaler
3 metadata:
4   name: test-hpa
5 spec:
6   scaleTargetRef:
7     apiVersion: apps/v1
8     kind: Deployment
9     name: test-deployment
10  minReplicas: 3
11  maxReplicas: 6
12  metrics:
13    - type: Resource
14      resource:
15        name: cpu
16        target:
17          type: Utilization
18          averageUtilization: 15
19  behavior:
20    scaleDown:
21      policies:
22        - type: Percent
23          value: 10
```

原来的40%要求优点高，我们尝试降低到15%，这样观察的效果更明显

效果如下：

```
^Croot@k8s-master:/home./test_add.sh
pod/pod-test1 created
pod/pod-test2 created
pod/pod-test3 created
pod/pod-test4 created
pod/pod-test5 created
pod/pod-test6 created
pod/pod-test7 created
pod/pod-test8 created
pod/pod-test9 created
pod/pod-test10 created
pod/pod-test11 created
pod/pod-test12 created
```

```
root@k8s-master:/home# kubectl get hpa test-hpa --watch
```

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
test-hpa	Deployment/test-deployment	0%/15%	3	6	3	59m
test-hpa	Deployment/test-deployment	0%/15%	3	6	3	59m
test-hpa	Deployment/test-deployment	1%/15%	3	6	3	60m
test-hpa	Deployment/test-deployment	10%/15%	3	6	3	60m
test-hpa	Deployment/test-deployment	18%/15%	3	6	3	60m
test-hpa	Deployment/test-deployment	20%/15%	3	6	4	60m
test-hpa	Deployment/test-deployment	20%/15%	3	6	4	61m
test-hpa	Deployment/test-deployment	21%/15%	3	6	4	61m
test-hpa	Deployment/test-deployment	21%/15%	3	6	4	61m
test-hpa	Deployment/test-deployment	20%/15%	3	6	4	61m

我们可以看到对应的负载逐渐变高，当CPU利用率达到20时，发生了自动扩容，replicas的数量从3变到了4。

我们同样编写删除脚本test\_delete.sh，内容如下：

```
1 for i in {1..80}
2 do
3     podname="pod-test${i}"
4     kubectl delete pod ${podname} --grace-period=0 --force
5 done
```

也就是删除80个创建的pod，我们接着运行如下的指令。

```
^Croot@k8s-master:/home./test_delete.sh
Warning: Immediate deletion does not wait for confirmation that the running resource has been terminated. The resource may continue to run on the cluster indefinitely.
pod "pod-test1" force deleted
Warning: Immediate deletion does not wait for confirmation that the running resource has been terminated. The resource may continue to run on the cluster indefinitely.
pod "pod-test2" force deleted
Warning: Immediate deletion does not wait for confirmation that the running resource has been terminated. The resource may continue to run on the cluster indefinitely.
pod "pod-test3" force deleted
Warning: Immediate deletion does not wait for confirmation that the running resource has been terminated. The resource may continue to run on the cluster indefinitely.
pod "pod-test4" force deleted
Warning: Immediate deletion does not wait for confirmation that the running resource has been terminated. The resource may continue to run on the cluster indefinitely.
pod "pod-test5" force deleted
Warning: Immediate deletion does not wait for confirmation that the running resource has been terminated. The resource may continue to run on the cluster indefinitely.
pod "pod-test6" force deleted
Warning: Immediate deletion does not wait for confirmation that the running resource has been terminated. The resource may continue to run on the cluster indefinitely.
pod "pod-test7" force deleted
Warning: Immediate deletion does not wait for confirmation that the running resource has been terminated. The resource may continue to run on the cluster indefinitely.
pod "pod-test8" force deleted
Warning: Immediate deletion does not wait for confirmation that the running resource has been terminated. The resource may continue to run on the cluster indefinitely.
pod "pod-test9" force deleted
```

删除80个pod。

```

root@k8s-master:/home# kubectl get hpa test-hpa --watch
NAME          REFERENCE                TARGETS  MINPODS  MAXPODS  REPLICAS  AGE
test-hpa      Deployment/test-deployment 16%/15%   3         6         4         68m
test-hpa      Deployment/test-deployment 15%/15%   3         6         4         68m
test-hpa      Deployment/test-deployment 8%/15%    3         6         4         68m
test-hpa      Deployment/test-deployment 11%/15%   3         6         4         68m
test-hpa      Deployment/test-deployment 10%/15%   3         6         4         69m
test-hpa      Deployment/test-deployment 2%/15%    3         6         4         69m
test-hpa      Deployment/test-deployment 0%/15%    3         6         4         69m
test-hpa      Deployment/test-deployment 0%/15%    3         6         4         70m
test-hpa      Deployment/test-deployment 0%/15%    3         6         4         70m
test-hpa      Deployment/test-deployment 0%/15%    3         6         4         72m
test-hpa      Deployment/test-deployment 0%/15%    3         6         4         72m
test-hpa      Deployment/test-deployment 0%/15%    3         6         4         73m
test-hpa      Deployment/test-deployment 0%/15%    3         6         3         73m
test-hpa      Deployment/test-deployment 0%/15%    3         6         3         74m
test-hpa      Deployment/test-deployment 0%/15%    3         6         3         74m

```

CPU利用率逐渐降低，最后降到了0，这个时候 hpa 缩容了，也就是replicas的数量重新由4变回了3。至此已经可以说明我们的hpa配置部署成功。